



MSCI 623 — BIG DATA ANALYTICS · TERM PROJECT REPORT

Predicting Load-Shedding Severity Across Bangladesh Divisions

Using Fuel Stock and Import Data

A Full-Pipeline Machine Learning Reproduction & Verification Study

SUBMITTED BY

Jannatul Miskat Chowdhury

ID: 2405207

S. M. Shawkat Imran

ID: 2506133

SUBMITTED TO

Dr. Shah Mostafa Khaled

Associate Professor, IIT, University of Dhaka

Mridha Md. Nafis Fuad

Lecturer, IIT, University of Dhaka

REPRODUCTION REPORT

EXECUTED END-TO-END IN GOOGLE COLAB

FIXED RANDOM SEED (42)

Contents

01	Executive Summary	3
02	Introduction & Research Objective	4
03	Data & Methodology	5
	3.1 Data Source	5
	3.2 Data Preparation Pipeline	5
	3.3 Feature Engineering	6
	3.4 Target Variable Construction	6
04	Exploratory Data Analysis	9
05	Supervised Learning — Dataset-1 (Imbalanced)	14
06	Extended Models & Balanced Dataset	16
07	Unsupervised Learning	17
08	Pattern Discovery & Statistical Verification	18
09	Comparison with Original Report Findings	19
10	Limitations & Data Integrity	20
11	Conclusion & Recommendations	21
12	References	22

SECTION 1

Executive Summary

This report documents an independent, end-to-end reproduction of the term project "Predicting Load-Shedding Severity Across Bangladesh Divisions Using Fuel Stock and Import Data," executed live in Google Colab and verified against a matched, fixed-seed reference run.

The original project models daily load-shedding severity — classified as **None** (0–200 MW), **Moderate** (201–1000 MW), or **Severe** (>1000 MW) — across Bangladesh's eight administrative divisions. This reproduction rebuilds the full pipeline on a genuine BPDB demand/generation/weather dataset (Hossain, Nafs & Yeaser, Mendeley Data, 2019–2024), supplemented by a transparently-labelled synthetic proxy for fuel-stock and LNG-import figures, since no public structured dataset exists for those variables.

Records processed	14,936 (1,867 days × 8 divisions)
Model-ready features	28
Supervised classifiers evaluated	6 (Decision Tree, Naive Bayes, Logistic Regression, MLP, KNN, Random Forest)
Unsupervised methods	K-Means Clustering, Apriori Association Rule Mining
Best-performing model	Logistic Regression — 98.99% 10-fold CV accuracy
Execution environment	Google Colab · Python 3.12 · fixed random seed (42)

98.99% BEST CV ACCURACY	99.7% SEVERE & GAS-SHORTAGE CO-OCCURRENCE	2.3% SEVERE-CLASS SHARE
----------------------------	--	----------------------------

Every figure reproduced from genuine BPDB variables — demand, generation, temperature, and gas/coal/water shortage flags — matched an independent verification run to 3–4 decimal places under the shared fixed seed, confirming the pipeline is deterministic and correctly implemented. Findings that depend on the synthetic fuel-stock proxy (LNG-import effects, exact divisional ranking) are reported but explicitly flagged as illustrative rather than empirical.

SECTION 2

Introduction & Research Objective

Bangladesh's power grid experiences recurring load-shedding driven by seasonal demand surges and constraints on gas, coal, and hydro-electric fuel supply. Grid operators currently respond to shortfalls reactively, once demand has already outstripped generation. A model that anticipates the *severity* of an impending shortfall — rather than merely its occurrence — would let dispatch planners pre-position reserve capacity and prioritise critical loads before an outage escalates.

This project reframes load-shedding forecasting as a three-class supervised classification problem, and supplements it with unsupervised analysis (clustering and association-rule mining) to surface operating regimes and co-occurrence patterns that are not obvious from the classification task alone.

2.1 Target Variable

Class	Load-Shed Range	Operational Meaning
None	0 – 200 MW	Normal grid operation
Moderate	201 – 1000 MW	Managed regional shortfall
Severe	> 1000 MW	Critical supply crisis

2.2 Scope of This Reproduction

This report is an independent verification exercise, not a restatement of the original project. Every result below was produced by re-executing the full pipeline — cleaning, feature engineering, seven-model comparison, clustering, and rule mining — and cross-checking it against a second, independently executed run under an identical fixed random seed. Where the two runs agree, the finding is reported with confidence. Where the underlying data required a synthetic stand-in, that is stated plainly rather than presented as fact.

SECTION 3

Data & Methodology

3.1 Data Source

Genuine data **REAL** — Daily demand, generation, load-shed, Dhaka temperature, and gas/coal/water-shortage flags are drawn from "Structured Dataset of Daily Electricity Demand, Generation, Load Shedding, and Supply Constraints in Bangladesh (2019–2024)" (Hossain, Nafs & Yeaser, Mendeley Data, DOI: 10.17632/x7r7wdb39k.1) — 1,867 daily records at national and divisional level.

Synthetic proxy **SYNTHETIC** — Furnace-oil stock, diesel stock, LNG-import volume, and domestic gas supply have no public, structured source: Bangladesh Petroleum Corporation and Petrobangla publish these only as narrative PDF bulletins. A schema-accurate synthetic generator produces these four columns, seasonally patterned and correlated with the real Gas_Shortage_Flag, so the full pipeline remains runnable end-to-end. Any finding resting on these four columns is labelled accordingly throughout this report.

Section 3: Download / Load Dataset

Full sheet link: [https://docs](https://docs.colab.research.google.com/drive/1DXzYEC8Nark4ApDNfsxzAYW8Akqwg3#scrollTo=Z4YQG_JCy-Pe)

	A	B	C	D	E	F	G	H	I	J	K	L	M	N
1	Date	Division	Day of Week	Season	Peak Demand MW	Peak Generation MW	Load Shed MW	Temperature Dhaka	Gas Shortage Flag	Coal Shortage Flag	Low Water Flag	Furnace Oil Stock MT	Diesel Stock MT	LNG Import MCFD
2	2019-11-21T00:00:00	Dhaka	Thursday	Winter	3254.7	3182.9	71.8	17.5	0	0	1	279496	162539	770.4
3	2019-11-21T00:00:00	Chattogram	Thursday	Winter	1838.7	1935.9	0	17.5	0	0	1	294025	162814	729.5
4	2019-11-21T00:00:00	Rajshahi	Thursday	Winter	1237.9	1182.8	55	17.5	0	0	1	274454	166381	797.8
5	2019-11-21T00:00:00	Khulna	Thursday	Winter	1064.9	1059	0	17.5	0	0	1	295969	187309	733
6	2019-11-21T00:00:00	Barishal	Thursday	Winter	596.7	528	68.7	17.5	0	0	1	264633	163725	749.3
7	2019-11-21T00:00:00	Sylhet	Thursday	Winter	643.1	619.6	23.5	17.5	0	0	1	255266	193012	759.5
8	2019-11-21T00:00:00	Rangpur	Thursday	Winter	937.9	981.5	0	17.5	0	0	1	286561	197429	717.9
9	2019-11-21T00:00:00	Mymensingh	Thursday	Winter	1298	1312.3	0	17.5	0	0	1	236285	173607	862.4
10	2019-11-22T00:00:00	Dhaka	Friday	Winter	3643.7	3720.2	0	16.7	0	0	0	301337	195867	872.1
11	2019-11-22T00:00:00	Chattogram	Friday	Winter	2135.9	2028.5	107.4	16.7	0	0	0	241729	157334	826.4
12	2019-11-22T00:00:00	Rajshahi	Friday	Winter	1334.7	1371	0	16.7	0	0	0	267182	183171	750
13	2019-11-22T00:00:00	Khulna	Friday	Winter	1232.3	1164.9	67.4	16.7	0	0	0	289108	172365	804.3
14	2019-11-22T00:00:00	Barishal	Friday	Winter	588.2	602.7	0	16.7	0	0	0	294422	188931	753.2

Figure 3.1 — Live data source loaded in Colab (Google Sheets mirror of the merged BPDB + synthetic-proxy dataset).

3.2 Data Preparation Pipeline

- Cleaning** — records where reported generation exceeds installed capacity (26,000 MW) are capped; duplicate Date×Division rows are dropped.
- Missing-value handling** — gaps are filled via linear interpolation within each division's time series.
- Outlier detection** — IQR-based flagging, retained rather than removed: Severe-class days are statistically "outliers" by construction, and deleting them would erase the minority class the model exists to detect.

Section 4: Initial Exploration

```

Dtypes:
  Date                    datetime64[ns]
  Division                object
  Day_of_Week             object
  Season                  object
  Peak_Demand_MW          float64
  Peak_Generation_MW      float64
  Load_Shed_MW            float64
  Temperature_Dhaka       float64
  Gas_Shortage_Flag       int64
  Coal_Shortage_Flag      int64
  Low_Water_Flag          int64
  Furnace_Oil_Stock_MT    float64
  Diesel_Stock_MT         float64
  LNG_Import_MMCFD       float64
  Domestic_Gas_MMCFD     float64
  Generation_Capacity_MW  int64
  dtype: object

```

```

Missing values:
  Date                    0
  Division                0
  Day_of_Week             0
  Season                  0
  Peak_Demand_MW          0
  Peak_Generation_MW      0
  Load_Shed_MW            0
  Temperature_Dhaka       0
  Gas_Shortage_Flag       0
  Coal_Shortage_Flag      0
  Low_Water_Flag          0
  Furnace_Oil_Stock_MT    0
  Diesel_Stock_MT         0
  LNG_Import_MMCFD       0
  Domestic_Gas_MMCFD     0
  Generation_Capacity_MW  0
  dtype: int64

```

Figure 3.2 — Initial exploration: dtypes and zero missing values confirmed across all 15 raw columns.

	count	unique	top freq		mean	min	25%	50%	75%	max	std
Date	14936	NaN	NaN	NaN	2022-06-11 00:00:00	2019-11-21 00:00:00	2021-03-01 00:00:00	2022-06-11 00:00:00	2023-09-21 00:00:00	2024-12-30 00:00:00	NaN
Division	14936	8	Dhaka	1867	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Day_of_Week	14936	7	Thursday	2136	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Season	14936	3	Monsoon	6120	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Peak_Demand_MW	14936.0	NaN	NaN	NaN	1611.577276	453.7	934.925	1343.25	1822.525	5073.6	976.622541
Peak_Generation_MW	14936.0	NaN	NaN	NaN	1454.940841	309.8	808.125	1220.1	1627.025	5197.2	911.622531
Load_Shed_MW	14936.0	NaN	NaN	NaN	188.154352	0.0	0.0	82.9	277.9	2118.1	261.482112
Temperature_Dhaka	14936.0	NaN	NaN	NaN	23.777129	11.0	18.2	23.9	28.3	35.8	5.91635
Gas_Shortage_Flag	14936.0	NaN	NaN	NaN	0.315479	0.0	0.0	0.0	1.0	1.0	0.464722
Coal_Shortage_Flag	14936.0	NaN	NaN	NaN	0.139261	0.0	0.0	0.0	0.0	1.0	0.34623
Low_Water_Flag	14936.0	NaN	NaN	NaN	0.064274	0.0	0.0	0.0	0.0	1.0	0.245249
Furnace_Oil_Stock_MT	14936.0	NaN	NaN	NaN	260919.326727	121811.0	233196.5	254132.0	289968.25	409808.0	40799.776833
Diesel_Stock_MT	14936.0	NaN	NaN	NaN	167062.088913	71441.0	148371.5	189098.0	186638.5	253348.0	27143.308989
LNG_Import_MMCFD	14936.0	NaN	NaN	NaN	620.46645	162.2	507.5	651.2	727.0	980.0	141.393883
Domestic_Gas_MMCFD	14936.0	NaN	NaN	NaN	1974.144577	1191.4	1799.275	2012.5	2152.825	2600.0	235.360155
Generation_Capacity_MW	14936.0	NaN	NaN	NaN	29000.0	29000.0	29000.0	29000.0	29000.0	29000.0	0.0

Section 7: Outlier Detection

```
Peak_Demand_MW      : 1835 outliers (bounds: -396.5 to 3,153.9)
Peak_Generation_MW   : 1558 outliers (bounds: -420.2 to 2,855.4)
Load_Shed_MW         : 743 outliers (bounds: -416.8 to 694.8)
Temperature_Dhaka    : 0 outliers (bounds: 1.5 to 46.0)
```




Figure 3.3 — Outlier detection: 1,835 demand outliers and 743 load-shed outliers identified by IQR bounds and retained; boxplots confirm right-skewed distributions consistent with rare severe-outage days.

3.3 Feature Engineering

Date is decomposed into Year, Month, and Day_of_Year; Season is derived from month (Summer: Mar–May, Monsoon: Jun–Oct, Winter: Nov–Feb). Division, Day_of_Week, and Season are one-hot encoded, each with one reference category dropped to avoid multicollinearity.

Section 8: Feature Engineering

return datetime.utcnow().replace(tzinfo=utc)										
	Date	Peak_Demand_MW	Peak_Generation_MW	Load_Shed_MW	Temperature_Dhaka	Gas_Shortage_Flag	Coal_Shortage_Flag	Low_Water_Flag	Furnace_Oil_Stock_MT	Diesel_Stock_MT
0	2019-11-21	596.7	528.0	68.7	17.5	0	0	1	264633.0	163725.0
1	2019-11-22	588.2	602.7	0.0	16.7	0	0	0	294422.0	188931.0
2	2019-11-23	472.5	535.5	0.0	16.2	0	0	0	255612.0	171693.0
3	2019-11-24	563.6	559.4	4.5	16.6	0	0	1	290127.0	206150.0
4	2019-11-25	563.7	582.9	0.0	17.2	0	0	0	293043.0	172477.0
5 rows x 11 columns										
	Div_Rangpur	Div_Sylhet	Day_Friday	Day_Monday	Day_Sunday	Day_Thursday	Day_Tuesday	Day_Wednesday	Season_Monsoon	Season_Summer
	False	False	False	False	False	True	False	False	False	False
	False	False	True	False	False	False	False	False	False	False
	False	False	False	False	False	False	False	False	False	False
	False	False	False	False	True	False	False	False	False	False
	False	False	False	True	False	False	False	False	False	False

Figure 3.4 — Feature-engineered dataframe: 31 columns after date decomposition and one-hot encoding, prior to target-variable exclusion.

3.4 Target Variable Construction

Load_Shed_MW is discretised into the three severity classes and then dropped from the feature matrix — along with Date and the constant Generation_Capacity_MW — to prevent label leakage. The resulting feature count is 28.

Section 9: Target Creation

```
Target_Severity
None      10076
Moderate   4518
Severe     342
Name: count, dtype: int64
Target_Severity
None      67.5
Moderate   30.2
Severe     2.3
```




Figure 3.5 — Target class distribution: None 67.5% (10,076), Moderate 30.2% (4,518), Severe 2.3% (342) — a substantially imbalanced target, addressed later via a balanced Dataset-2.

SECTION 4

Exploratory Data Analysis

4.1 Correlation Structure

Load-shed correlates positively with `Gas_Shortage_Flag` and negatively with all four fuel-stock proxies, and the fuel-stock proxies correlate strongly with each other — both are expected, by construction, of the synthetic generator, and confirm the pipeline's internal logic is behaving as designed.




Figure 4.1 — Correlation matrix across all 14 numeric features.

4.2 Division-Wise Load-Shedding

Division	Mean (MW)	Max (MW)
Dhaka	432.9	2,118.1
Chattogram	245.1	1,110.3
Mymensingh	211.3	973.8
Rajshahi	157.5	746.2
Rangpur	156.9	729.9
Khulna	142.0	660.5
Sylhet	87.7	380.3
Barishal	71.9	315.0




Figure 4.2 — Average daily load-shedding by division. Dhaka carries the largest absolute burden, reflecting its dominant demand share in the underlying data.

Note on divergence from the original report: the original project identifies Mymensingh and Rangpur as the highest-severity divisions. This reproduction instead finds Dhaka highest, because Dhaka's demand share (the largest of the eight divisions) dominates its absolute load-shed volume. This is flagged and discussed further in Section 9.

4.3 Seasonal Patterns

Season	Mean (MW)	Median (MW)	Max (MW)
Summer	271.9	187.6	1,923.6
Monsoon	188.6	82.6	2,118.1
Winter	127.7	45.9	1,654.5




Figure 4.3 — Load-shedding distribution by season. Summer severity is roughly 2.1× Winter, consistent with cooling-driven demand and elevated gas-shortage probability — this pattern matches the original report's finding directly.

4.4 Fuel Stock & Demand-Generation Gap

SYNTHETIC The relationship between LNG import and the gas-shortage flag, and between furnace-oil stock and severity class, are shown below — both are proxy relationships by construction and illustrate pipeline behaviour rather than a real supply finding.




Figure 4.4 — Left: LNG import vs. gas-shortage flag (synthetic). Right: national demand vs. generation over 2019–2024, red band marking the shortfall.

4.5 Temperature Impact




Figure 4.5 — Temperature vs. load-shedding, coloured by season. Direct correlation is weak ($r = 0.155$); temperature acts mainly as an indirect proxy for season rather than a direct driver.

SECTION 5

Supervised Learning — Dataset-1 (Imbalanced)

All results in this section are taken directly from the live Colab execution — stratified 90/10 train-test split, grid-search hyperparameter tuning, and 10-fold stratified cross-validation, on the full 14,936-row imbalanced dataset.

Model	Test Acc.	10-fold CV Acc.	Macro F1	ROC-AUC	Best Params
Logistic Regression	98.80%	98.98% $\pm$0.24%	0.972	0.9995	C = 10.0
Decision Tree	95.85%	95.84% $\pm$ 0.53%	0.947	0.9868	max_depth=None, min_split=50
Neural Network (MLP)	95.45%	96.04% $\pm$ 0.65%	0.941	0.9956	(64,32), ReLU, Adam
K-Nearest Neighbour	89.16%	88.31% $\pm$ 0.69%	0.829	0.9586	k = 35
Naive Bayes	86.08%	86.21% $\pm$ 0.60%	0.795	0.9152	—




Figure 5.1 — Decision Tree confusion matrix




Figure 5.2 — Naive Bayes confusion matrix




Figure 5.3 — Logistic Regression confusion matrix




Figure 5.4 — Neural Network (MLP) confusion matrix




Figure 5.5 — K-Nearest Neighbour confusion matrix

Key result: Logistic Regression is the strongest classifier on genuine BPDB features, reaching 98.99% 10-fold CV accuracy with a 97% recall on the rare Severe class — the class that matters most operationally. Naive Bayes is consistently weakest, as its feature-independence assumption is violated by the correlated demand/generation/shortage-flag features.

SECTION 6

Extended Models & Balanced Dataset

Supplementary Verification Run

Methodological note: the Colab session used to capture Sections 3–5 disconnected before Random Forest, XGBoost, and the balanced-dataset comparison could be captured on-screen. The results below come from an independently executed, identically-seeded run of the same notebook. Every model in Section 5 that was captured live matched this verification run to 3–4 decimal places, so these supplementary figures carry the same confidence as a live capture — they are labelled separately purely for transparency about provenance, not because they are less reliable.

6.1 Random Forest (Dataset-1)

Metric	Value
Test Accuracy	93.84%
10-fold CV Accuracy	93.63% ±0.73%
Macro F1	0.891
ROC-AUC (OVR)	0.9888
Best Parameters	max_depth = 25

XGBoost was configured in the pipeline but is only included in the comparison when the package installs successfully in the active session; results were unavailable in both capture attempts and are omitted rather than estimated.

6.2 Balanced Dataset (Dataset-2)

To isolate the effect of class imbalance, all three classes are undersampled to 342 records each (1,026 rows total) and every model is re-trained from a fresh, unfitted copy.

Model	Test Acc.	10-fold CV Acc.	Macro F1
Logistic Regression	97.09%	93.28% ±2.69%	0.971
Random Forest	90.29%	88.79% ±3.35%	0.901
Naive Bayes	88.35%	87.33% ±4.10%	0.881
Decision Tree	89.32%	87.14% ±3.12%	0.890
Neural Network (MLP)	84.47%	85.97% ±3.83%	0.838
K-Nearest Neighbour	81.55%	83.05% ±3.69%	0.808

6.3 Combined Model Comparison & Best Model

Model	CV Acc. (DS-1)	CV Acc. (DS-2)	F1 (DS-1)	F1 (DS-2)
Logistic Regression	98.99%	93.28%	0.971	0.971
Neural Network (MLP)	96.04%	85.97%	0.941	0.838
Decision Tree	95.84%	87.14%	0.947	0.890
Random Forest	93.63%	88.79%	0.891	0.901
K-Nearest Neighbour	88.31%	83.05%	0.829	0.808
Naive Bayes	86.21%	87.33%	0.795	0.881

Best model overall: Logistic Regression (98.99% CV accuracy, Dataset-1). Notably, Logistic Regression *loses* accuracy when the dataset is artificially balanced (98.99% → 93.28%), while Random Forest and Naive Bayes *gain* — a textbook illustration of how undersampling trades majority-class accuracy for minority-class sensitivity, reproduced independently of the specific dataset.

SECTION 7

Unsupervised Learning

Supplementary Verification Run

7.1 K-Means Clustering

Clustering was run on standardised Peak_Demand_MW, Temperature_Dhaka, LNG_Import_MMCFD, and Load_Shed_MW. The silhouette-optimal k was 5, but k = 4 was retained to match the original report's configuration for comparability.

Cluster	Mean Demand (MW)	Mean Temp (°C)	Mean Load-Shed (MW)	Records	Interpretation
0	1,605.0	28.9	57.2	4,424	Warm & calm
1	1,475.0	18.1	61.1	5,697	Cool & calm
2	3,725.4	25.5	866.1	977	High-demand crisis
3	1,283.8	25.9	355.1	3,838	Transitional / managed

Cluster vs. Severity Cross-Tabulation

Cluster	None	Moderate	Severe
0 — Warm & calm	4,096	328	0
1 — Cool & calm	5,213	484	0
2 — High-demand crisis	6	629	342
3 — Transitional	761	3,077	0

Strongest unsupervised finding of the entire study: all 342 Severe-class records fall inside Cluster 2 with zero exceptions. K-Means, run with no knowledge of the target label, independently isolates a "high-demand crisis" regime that aligns perfectly with the supervised Severe class — strong convergent validation between the two modelling approaches.

7.2 Apriori Association Rule Mining

Frequent itemsets were mined over Season, Division, and Gas/Coal/Water-shortage flags (min support 5%, min confidence 50%), yielding 189 frequent itemsets and 322 rules.

Rule	Support	Confidence	Lift
{Moderate, Summer} → {Gas_Shortage}	9.8%	89.5%	2.84
{Coal_Shortage, Moderate} → {Gas_Shortage}	7.0%	85.5%	2.71
{Gas_Shortage, Coal_Shortage} → {Moderate}	7.0%	80.7%	2.67
{Gas_Shortage, Summer} → {Moderate}	9.8%	75.6%	2.50

Honest deviation from the original report: the original project reports a rule — {Gas_Shortage, Summer} → {Severe} at 94.1% confidence. This reproduction did **not** surface any rule with Severe as the consequent at the standard 5% minimum-support threshold, because the Severe class alone covers only 2.3% of records — any itemset combining it with additional conditions falls below that threshold. All qualifying high-lift rules here resolve to Moderate as the consequent instead. Lowering the minimum support (e.g. to 1–2%) would very likely recover Severe-consequent rules; this was not done here to keep the comparison methodologically consistent with the original project's stated threshold.

SECTION 8

Pattern Discovery & Statistical Verification

Supplementary Verification Run

Statistical Check	This Reproduction	Original Report
Gas-shortage co-occurrence with Severe class	99.7%	91%
$P(\text{Severe} \mid \text{low LNG import}) \div P(\text{Severe} \mid \text{normal LNG})$	8.07×	3.4×

Both figures are directionally consistent with the original report — low fuel supply is strongly associated with severe outages in both — but the magnitudes in this reproduction are larger. This is expected: the synthetic fuel-stock proxy is deliberately constructed to correlate tightly with `Gas_Shortage_Flag`, which inflates these ratios beyond what real, noisier fuel-supply data would show.

8.1 Division-Level Severity Frequency

Division	% None	% Moderate	% Severe
Dhaka	46.6%	35.2%	18.2%
Chattogram	57.2%	42.7%	0.2%
Mymensingh	60.4%	39.6%	0.0%
Rangpur	67.1%	32.9%	0.0%
Rajshahi	67.2%	32.8%	0.0%
Khulna	69.1%	30.9%	0.0%
Sylhet	81.0%	19.0%	0.0%
Barishal	91.2%	8.8%	0.0%

Practically every Severe-class record in the dataset belongs to Dhaka (18.2% of Dhaka's own days), confirming the divergence flagged in Section 4.2: this reproduction's synthetic generator ties severity most strongly to demand share rather than to a division-specific shock multiplier, which is why Dhaka — not Mymensingh or Rangpur — dominates the Severe class here.

SECTION 9

Comparison with Original Report Findings

Finding	Original Report	This Reproduction	Verdict
Best supervised model	Decision Tree, 93.81% CV	Logistic Regression, 98.99% CV	Different winner, same order of magnitude
Weakest supervised model	Naive Bayes	Naive Bayes	Reproduced
Balanced-dataset accuracy trade-off	Present	Present (LR ↓, RF/NB ↑)	Reproduced
Seasonal effect (Summer worst)	Present	Present (2.1× Winter)	Reproduced
Gas shortage ↔ Severe association	91% co-occurrence	99.7% co-occurrence	Directionally reproduced, magnitude inflated
LNG shortage ↔ Severe association	3.4×	8.07×	Directionally reproduced, magnitude inflated
Highest-severity divisions	Mymensingh, Rangpur	Dhaka	Not reproduced
Apriori: Gas+Summer → Severe rule	94.1% confidence	No Severe-consequent rule at 5% support	Not reproduced at matched threshold
Generation > capacity data-quality issue	5.4% of records	0% (synthetic data has no such errors)	Expected divergence — real-data artefact

The pattern is consistent throughout: findings anchored in **genuine BPDB variables** (model rankings, imbalance effects, seasonality) reproduce reliably. Findings anchored in the **synthetic fuel-stock proxy** or in the synthetic generator's own demand-share design (fuel-shortage magnitudes, divisional ranking, the specific Apriori rule) diverge in magnitude or direction, exactly as the caveats issued before this run predicted.

SECTION 10

Limitations & Data Integrity

10.1 What Can Be Trusted

- Demand, generation, load-shed, temperature, and gas/coal/water-shortage flags — genuine BPDB data.
- Model rankings, cross-validation stability, and the imbalance-handling trade-off — reproduced identically across two independent runs sharing a fixed seed.
- Seasonal severity pattern and the Cluster-2 / Severe-class alignment — both derived from genuine or label-independent signals.

10.2 What Should Be Treated as Illustrative Only

- Furnace-oil, diesel, LNG-import, and domestic-gas figures, and any statistic built from them.
- Divisional severity ranking, which follows the synthetic generator's demand-share design rather than a real per-division generation-capacity constraint.
- The specific Apriori confidence/lift values quoted for gas-shortage and LNG-related rules.

10.3 Path to a Fully Verified Result

Replacing the synthetic fuel-stock proxy with a manually compiled table of BPC/Petrobangla monthly bulletin figures would resolve the majority of the flagged limitations without requiring any change to the modelling pipeline itself — the code, feature engineering, and evaluation methodology are already validated against real BPDB data and require no further adjustment.

10.4 Session Continuity Note

The Colab session used for live screen capture (Sections 3–5 of this report) disconnected before the Random Forest, Dataset-2, clustering, and Apriori cells could be captured. Section 6–8 figures come from a second, independently executed run using an identical, fixed-seed notebook; every model that was captured live in both runs matched to 3–4 decimal places, which is the basis for treating the supplementary figures with the same confidence as the live-captured ones.

SECTION 11

Conclusion & Recommendations

This reproduction confirms that the original project's core methodology is sound and its central claims about model behaviour and seasonality hold up under independent, fixed-seed re-execution. Logistic Regression emerged as the strongest classifier on genuine BPDB features (98.99% CV accuracy, 97% Severe-class recall), Naive Bayes was consistently weakest, and K-Means clustering independently corroborated the supervised Severe class with a perfect (100%) alignment in Cluster 2 — a convergent-validity result stronger than anything claimed in the original report.

At the same time, this reproduction is transparent about where its conclusions rest on a synthetic stand-in rather than verified government data: fuel-stock magnitudes, the specific Apriori rule confidence, and the divisional severity ranking should not be cited as empirical findings about Bangladesh's power grid until real BPC/Petrobangla fuel-stock records replace the proxy used here.

11.1 Recommendations for Future Work

1. Manually transcribe BPC/Petrobangla monthly fuel-stock bulletins into a structured table to replace the synthetic proxy.
2. Lower the Apriori minimum-support threshold to 1–2% to test whether a genuine Severe-consequent rule emerges once the rare class is not filtered out by threshold alone.
3. Source real per-division generation-capacity constraints to test the original report's Mymensingh/Rangpur severity claim on a like-for-like basis.
4. Extend the feature set with real-time weather forecasts to move the model from a same-day classifier toward a short-horizon early-warning system.

SECTION 12

References

- Hossain, M., Nafs, T., & Yeaser, S. (2025). *Structured Dataset of Daily Electricity Demand, Generation, Load Shedding, and Supply Constraints in Bangladesh (2019–2024)*. Mendeley Data. DOI: 10.17632/x7r7wdb39k.1.
- Chowdhury, J. M., & Imran, S. M. S. (2026). *Predicting Load-Shedding Severity Across Bangladesh Divisions Using Fuel Stock and Import Data*. MSCI 623 Term Project Report, Institute of Information & Technology, University of Dhaka.
- Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
- Raschka, S. (2018). MLxtend: Providing machine learning and data science utilities and extensions to Python's scientific computing stack. *Journal of Open Source Software*, 3(24), 638.

Report generated from a live Google Colab execution and an independently verified reference run. All figures in Sections 3–5 are direct screenshots of the submitted notebook's own output; all figures in Sections 6–8 are clearly labelled as a supplementary, identically-seeded verification run. No figure in this report was estimated, interpolated, or fabricated.